

Rust Type Layout

Size, Alignment, Endianess, and
Representations
VS
Input, Output, and Foreign Functions

Objective: Binary Compatible

- I want to be able to exchange data structures between
- multiple instances of my application running on different processor architectures like x86-64, ARM, Power, RISC-V, s390x
 - <https://doc.rust-lang.org/rustc/platform-support.html>, and between
- my application and other applications – possibly written in other programming languages.

Size depends on the Platform

Fixed-Size Types

- Some primitive types have fixed sizes.

Type	Size in 8-bit bytes	Description
i8	1	8-bit signed two's complement binary integer
u8	1	8-bit unsigned binary integer
i16	2	16-bit signed two's complement binary integer
u16	2	16-bit unsigned binary integer
i32	4	32-bit signed two's complement binary integer
u32	4	32-bit unsigned binary integer
i64	8	64-bit signed two's complement binary integer
u64	8	64-bit unsigned binary integer
i128	16	128-bit signed two's complement binary integer
u128	16	128-bit unsigned binary integer
f32	4	32-bit binary IEEE 754-2008 floating-point number
f64	8	64-bit binary IEEE 754-2008 floating-point number
()	0	Unit type
bool	1	Boolean true or false
char	4	UTF-32 Unicode character

Varying-Size Types

- For other primitive types their sizes depend on the platform.

Type	Size in 8-bit bytes	Description
isize	?	signed binary integer size
usize	?	unsigned binary integer size
fn	?	function pointer
*T	?	pointer
&T	?	reference
Box<T>	?	box
Option<T>	?	option
Option<Box<T>>	?	optional box

Size Conclusion

- Use only fixed-size types in external data structures.
- Get their sizes via `std::mem::size_of<T>()`.
- While `size_of<T>()` is a constant function, for varying-size types it can return different results on different architectures and theoretically even between different Rust releases.

Alignment depends on the Platform

Alignment depends on Platform

- The alignment depends entirely on the platform.
- The function `std::mem::align_of<T>()` returns the minimum required alignment in structs for the type `T`.
 - This can differ from the preferred alignment.
 - There is no function in `std::mem` to obtain this preferred alignment.
 - https://doc.rust-lang.org/std/mem/fn.align_of.html
- For example, on the x86-64 architecture the 128-bit binary integer types are aligned on 64 bits.
- `std::alloc::Layout` supports alignment on powers of 2.

Alignment Conclusions

- You can determine the minimum required alignment simply and quickly by running `std::mem::align_of<T>()` for each primitive type on each platform that is relevant for you.
- `std::mem::align_of<T>() <= std::mem::size_of<T>()` holds true because of the definition of `size_of`.
 - Therefore it is safe to align types on their size.
- To determine the preferred alignment, you have to read the platform specification.

Endianess depends on the Platform

Endianness depends on Platform

- Integer primitive types longer than 1 byte can be big or little endian. So far Rust does not list any mixed-endian platforms.
- The methods `from_be(x)`, `from_le(x)`, `to_be()`, and `to_le()` make the endianness explicit, but work on the same integer type only.
- The methods `from_be_bytes(x)`, `from_le_bytes(x)`, `to_be_bytes()`, and `to_le_bytes()` are defined for byte arrays (`[u8; n]`) only.
- There are no endianness functions that work on byte slices (`[u8]`).

Endianness Conclusions

- Explicitly specify the endianness of external integer fields via the methods `from_be(x)`, `from_le(x)`, `to_be()`, and `to_le()`.
- Explicitly specify the endianness of external integer fields in byte buffers via the methods `from_be_bytes(bytes)`, `from_le_bytes(bytes)`, `to_be_bytes()`, and `to_le_bytes()`.

Representations

Representations

- The standard Rust representation guarantees only soundness regarding the type layout of user defined composite types (structs, enums, and unions).
- The C representation guarantees platform-independent ordering and alignment, but the size of enums is excepted.
- The packed representation has defined alignments, but can create unaligned items.
- See <https://doc.rust-lang.org/reference/type-layout.html>

Representation Conclusions

- Use the C representation `repr(c)` to enforce the ordering and alignment of fields within structures.
- Explicitly specify reserved fields instead of relying on the insertion of padding bytes to properly align fields.

Input and Output of Bytes only

Input and Output Traits

- `std::io::Read` reads into byte slices and byte vectors.
- `std::io::Write` writes from byte slices.
- `std::fs::File` implements Read and Write.
- `std::net::TcpStream` implements Read and Write.
- Byte slices have an alignment of 1.

Input and Output Conclusions

- There seem to be no safe interfaces in `std::io` for reading into or writing from slices of primitive types other than `u8`.
- Safely reading other primitive types requires constructing them from bytes after reading these.
- Safely writing other primitive types requires converting them to bytes before writing these.

Foreign Function Interface

External Functions

- Calls from Rust to non-Rust call functions declared in "extern" blocks with the specified ABI (e.g. "C").
- Functions in "extern" blocks are always "unsafe", because the Rust compiler could not check them.
- Calls from non-Rust to Rust call "extern" functions with the specified ABI.
- Extern functions are safe by default, but require the attribute `#[no_mangle]` to get stable external symbols.
- See <https://doc.rust-lang.org/book/ch19-01-unsafe-rust.html>

External Function Conclusion

- Non-Rust code is always unsafe from the perspective of Rust.
- You can avoid unsafe code by using IO instead of calls.
- IO comes with the overhead of system calls to the OS kernel.
- Also the OS kernel is potentially unsafe, unless the kernel is implemented in Rust as well, like in the Redox OS.
- On architectures with memory protection you can isolate unsafe code to separate processes.

End of Story

Comments or Questions?